



NATIONAL AND KAPODISTRIAN UNIVERSITY OF ATHENS

**FACULTY OF EXACT SCIENCES
DEPARTMENT OF INFORMATICS AND TELECOMMUNICATIONS**

BACHELOR THESIS

**RHEA: A Reactive, Heterogeneous, Extensible and Abstract
Framework for Dataflow Programming**

Orestis Melkonian

**Supervisors: Panos Rondogiannis, Professor EKPA
Angelos Charalambidis, Researcher NCSR**

ATHENS

APRIL 2016

BACHELOR THESIS

RHEA: A Reactive, Heterogeneous, Extensible and Abstract Framework for Dataflow
Programming

Orestis Melkonian

A.M.: 1115201000128

SUPERVISORS: **Panos Rondogiannis**, Professor EKPA
Angelos Charalambidis, Researcher NCSR

ABSTRACT

Summary here

SUBJECT AREA: Programming Languages

KEYWORDS: dataflow, stream, frp, distributed systems, model-driven architecture
program synthesis declarative languages node placement

”τὰ ὄντα ἰέναι τε πάντα καὶ μένειν ουδέν”
(all entities move and nothing remains still)
- Heraclitus

ACKNOWLEDGEMENTS

I would like to thank Angelos Charalambidis for his immensely helpful supervision and guidance throughout the whole period of 6 months that I was present in NCSR.

I would also like to thank Professor Panos Randoiannis for being a major influence in my current research interests through the undergraduate courses "Theory of Computation" and "Principles of Programming Languages", where I was introduced to a much more declarative way of programming and started appreciating mathematical sensibility in computing.

CONTENTS

PROLOGUE	10
1. INTRODUCTION	11
1.1 Main concept	11
1.2 Motivation	11
1.2.1 Declarative languages	11
1.2.2 Data versus Computation	11
1.2.3 Dataflows in Robotics	12
1.2.4 Dataflows in Big Data	12
1.3 Outline	12
2. BACKGROUND	13
2.1 The dataflow computational model	13
2.2 Functional reactive programming	13
2.3 Publish-Subscribe model	14
2.4 ROS: Robot Operating System	15
2.5 Internet of things	15
2.6 The Strategy design pattern	16
3. APPROACH	17
3.1 Reactive	17
3.2 Heterogeneous	18
3.3 Extensible	18
3.4 Abstract	18
4. IMPLEMENTATION	20
4.1 Notifications	20
4.2 The Reactive Streams Standard	20
4.3 Software structure	20

4.4	Internal representation	21
4.5	Stream variables	21
4.6	Stream operators	22
4.6.1	Creation	22
4.6.2	Combining	22
4.6.3	Filtering	22
4.6.4	Conditional	22
4.6.5	Transformational	22
4.6.6	Feedback	22
4.6.7	Error-handling	22
4.6.8	Backpressure	22
4.6.9	Utility	22
4.7	Evaluation	22
4.8	Distribution	22
4.8.1	Hazelcast	22
4.8.2	Machine configuration	22
4.9	Optimization	22
4.9.1	Heuristic-based graph transformations	22
4.9.2	Network-aware node placement	22
4.10	Serialization	22
4.10.1	The Serializer Interface	22
4.10.2	JsonIO	22
5.	APPLICATIONS	23
5.1	Hamming numbers	23
5.2	Camera surveillance	23
5.3	Robot control panel	23
5.4	Robot hospital guide	23
6.	RELATED WORK	24
6.1	Big Data	24
6.1.1	GoogleDataflow	24
6.1.2	TensorFlow	24
6.1.3	Akka	24

6.1.4	dispel4py	24
6.1.5	Spark	24
6.1.6	Naiad	24
6.2	Robotics	24
6.2.1	Flowstone	24
6.2.2	Yampa	24
6.2.3	roshask	24
6.3	Internet of Things	24
6.3.1	NodeRed	24
6.3.2	NoFlo	24
7.	FUTURE WORK	25
7.1	More evaluation strategies	25
7.2	Dynamic reconfiguration	25
7.3	Advanced network profiling	25
7.4	DSLs for framework extension	25
7.5	Advanced fault-tolerance	25
7.6	Integration with other technologies	25
7.7	Stream reasoning	25
8.	CONCLUSIONS	26
	ABBREVIATIONS, INITIALS AND ACRONYMS	27

LIST OF FIGURES

Figure 1:	Natural numbers	13
Figure 2:	PubSub typical layout	15
Figure 3:	Strategy design pattern	16
Figure 4:	Reactive properties	17
Figure 5:	Heterogeneity pipeline	18
Figure 6:	Simple example	21
Figure 7:	Split example	22

PROLOGUE

This bachelor thesis is a continuation of my internship at the National Centre for Scientific Research "Demokritos", particularly in the Software and Knowledge Engineering Laboratory (SKEL).

The main task I was assigned was the implementation of a framework for robot programming using a dataflow approach. During that internship, I came to realize that my work could be easily generalized to cover a much broader application area than just robot software.

The name of the framework stems from the ancient Greek Titaness Rhea(*Ρέα*), daughter of earth goddess Gaia(*Γαία*) and sky god Uranus(*Ουρανός*) and etymologically derives from the verb *ρέω*(to flow).

1. INTRODUCTION

1.1 Main concept

My main contribution is the design and implementation of a framework for dataflow programming to be deployed anywhere, ranging from low-performance robots and sensors to clusters of computer and even the Cloud.

The main idea is to provide the programmer with a different execution model, the dataflow model, which allows for a more abstract way of thinking and has the advantage of exposing opportunities for parallelism (amongst CPU cores) and distribution (amongst computational machines) that the "intelligent" underlying system can automatically realize. Therefore, the programmer will be able to gain good performance and utilization of the available computational resources, while at the same time reducing development time and cost and maintaining a much cleaner and easier-to-refactor software system.

1.2 Motivation

1.2.1 Declarative languages

Software is becoming increasingly more complex each year, as computing capabilities are strengthened and user needs become more and more demanding. Thus the need for higher abstraction becomes mandatory, as it provides a more structured, easier to debug and maintainable way of developing software. In other words, abstraction in computer science acts as a mean to overcome complexity.

In programming languages, the level of the aforementioned abstraction is measured regarding the amount of low-level details a programmer has to specify. Therefore, languages can be divided in two categories: the imperative ones, which specify what needs to be done and how to do it, and the declarative ones, which only specify what needs to be done and rely on the underlying compiler/interpreter to produce the exact commands that will realize the desired behaviour. The most well-known declarative programming paradigms are functional and logic programming, each providing higher abstraction in different aspects. My approach was greatly influenced by the functional paradigm.

1.2.2 Data versus Computation

A common problem in heterogeneous systems is that different representations of the same entities/data-types coexist in the same software and, as a consequence, pure computational tasks are intermingled with data-converting tasks. This makes the code less readable and harder to maintain and understand. In the dataflow execution model, there is a clear separation of these two aspects as data (edges) are completely decoupled from computation (nodes). This motivation is strengthened even more, when cross-machine communication is included, and apart from converting data from one representation to another, serialization(i.e. conversion to bytes) is also mandatory.

1.2.3 Dataflows in Robotics

In control theory, which is the main background theory used in robotics, most architectures and/or algorithms are represented as dataflow diagrams for the sake of clarity and intuition. Translating these diagrams into common "imperative" software is not an easy task and is usually the source of bugs. Thus, having a dataflow execution model will nullify the need for such a translation.

Moreover, robotics typically involve several different robotic systems, whose combination is even more challenging. If each individual system is represented as a dataflow graph, composing them together is as trivial as connecting inputs with outputs, which is not the case in a traditional architecture, which is not component-based.

1.2.4 Dataflows in Big Data

Another reason for following a dataflow approach is the attention that it recently has drawn in the Big Data field. As data size is growing exponentially and distribution is not a luxury but a necessity, a more scalable and decentralized architecture was destined to be examined in more depth. As we will discuss in the *Related Work* chapter, there are many recent frameworks that became popular for the scalability due to the fact that they rely on a dataflow approach.

1.3 Outline

There are nine chapters which compose this thesis: *Introduction*, *Background*, *Approach*, *Implementation*, *Applications*, *Related Work*, *Future Work*, *Conclusions*.

Background introduces the reader to basic background knowledge, necessary for complete understanding.

Approach presents the main characteristics of my approach.

Implementation gives a more detailed specification of the framework.

Applications present some use-cases, ranging from general mathematical problems to real-life robot scenarios.

Related Work discusses relevant concepts and technologies, which influenced major decisions concerning the design and implementation of the framework.

Future Work suggests some interesting topics for future research, whose embedding in the framework is meaningful.

Conclusions sums up.

2. BACKGROUND

2.1 The dataflow computational model

The increased interest in parallelism during the 70's gave rise to the dataflow execution model, which is an alternative to the classical "von-Neumann" model. In the dataflow model, everything is represented in a dataflow graph, where nodes are independent computational units and edges are communication channels between these units. A node/unit is fired immediately when its required inputs are available and therefore no explicit control commands are needed for execution. Figure 1 shows a dataflow graph enumerating the set \mathbb{N} of natural numbers.




Figure 1: Natural numbers

The main advantage of the dataflow model is its implicit parallelism, deriving from the fact that the computational units are totally independent and therefore can be executed in parallel. The communication can either be an in-memory data storage or even a TCP connection across the network. Its great flexibility and composability make it a good candidate for the underlying architecture of a framework with a high level of abstraction.

2.2 Functional reactive programming

A relatively recent model of programming is Functional Reactive Programming (FRP), which provides a conceptual framework for implementing reactive (i.e. time-varying and responding to external stimuli) behaviour in *hybrid systems* (i.e. containing both continuous and discrete components), such as robots, in functional programming languages. It first appeared as a composable library for graphic animations [1], but quickly evolved into a generic paradigm [2, 3, 4]. Moreover, extensive research has investigated FRP as a framework for robotics [5, 6].

Although appealing at first, FRP was not appropriate for systems with real-time constraints, due to uncontrollable time- and space- leaks [7]. The solution was a generalization of monads called *arrows* [8], which provided the necessary guarantees that the aforementioned common errors do not occur. Let's see the example of calculating a robot's x-coordinate. Here is the mathematical formula drawn from control theory:

$$x = 1/2 \int (vr + vl) \cos \theta$$

Below is the FRP code corresponding to the formula above:

```
x = let
  v = (vrSF &&& vlSF) >>> lift (+)
  t = thetaSF >>> arr cos
  in (v &&& t) >>> lift (*) >>> integral >>> lift (/2)
```

As the above may seem counter-intuitive and difficult to understand, a new notation was conceived, notably the *arrow notation* [9]:

```
x = proc inp -> do
  vr <- vrSF -< inp
  vl <- vlSF -< inp
  theta <- thetaSF -< inp
  i <- integral-< (vr+vl) * cos(theta)
returnA -< (i/2)
```

The main advantages of FRP are its close correspondence to mathematics, which make it an ideal framework for modelling real-time systems, and its concise representation of time-varying values via *signals*.

For a more detailed view of the FRP idiom and all of its variants please see the cited literature.

2.3 Publish-Subscribe model

Publish/Subscribe(PubSub) is a communication paradigm that became popular due to the loose coupling of its components, suited for the most recent large-scale distributed applications.

There is no point-to-point communication and no synchronization. *Publishers* advertise messages of a given type to a specific message class or *topic* that is identified by a *key-word*, whereas *subscribers* listen on a specific *topic* without any knowledge of who the publishers are. The component responsible for relaying the messages between machines and/or processes and finding the cheaper dissemination method is called *message broker*. Figure 2 illustrates an abstract representation of the PubSub model.




Figure 2: PubSub typical layout

2.4 ROS: Robot Operating System

ROS is an open-source middleware for robot software, which emphasizes large-scale integrative robotics research [10]. It provides a *thin* communication layer between heterogeneous computers, from robots to mainframes and it has been widely adopted by the research community around the world, due to its flexibility and maximal support of reusability through packaging and composability. It provides a nice solution to the development complexity introduced by complex robot applications that consist of several modules and require different device drivers for each individual robot.

It follows a peer-to-peer network topology, implemented using a topic-based PubSub messaging protocol and its architecture reflects many sound design principles. Another great property of ROS is that it is language-agnostic, meaning that only a minimal specification language for message transaction has been defined, so contributors are free to implement small-size clients in different programming languages, with *roscpp* and *rospy* being the most widely used ones.

A typical development scenario is to write several *nodes*, that subscribe to some topics and, after doing some computation, publish their results on other topics. The main architectural issue here is that subscribing is realized through asynchronous callback functions, so complicated schemes easily lead to unstructured code, which obviously lead to unreadable and hard-to-maintain code. My approach gives a solution to the aforementioned problem.

2.5 Internet of things

The birth of the Internet gave rise to a concept called *Internet of Things* (IoT), which is essentially the ability of many heterogeneous devices, ranging from low-cost sensors to vehicles with embedded electronics, to collect data and exchange it amongst themselves using the Internet. This gave rise to smart grids, smart homes and eventually smart cities.

The development of such systems though, due to their heterogeneity, is rather complex and costly. Typical software architectures were not meant to be used in such environments and therefore new tools and concepts needed to be invented.

Fortunately, the dataflow model seems to be rather fitting for these scenarios, as every node in the graph is completely independent, and consequently can be any *"thing"*. This useful property of the model makes it a good architectural choice for such applications.

The only thing to consider is how these things will communicate in a standard way, so as to be able to add new types of *things* and integrate it in an effortless way to an existing dataflow network of *things*.

2.6 The Strategy design pattern

In software engineering, and especially in object-oriented programming(OOP), a *design pattern* is a general repeatable solution to a commonly occurring problem in software design [11].

One such solution, the *Strategy* design pattern is used when a particular algorithm can be implemented by a variety of behaviours/classes [11]. In such a case, a good idea is to isolate the algorithm in a separate *interface* and allow the system to select the appropriate instantiating classes at runtime.

Figure 3 illustrates the basic UML diagram of the strategy design pattern.




Figure 3: Strategy design pattern

3. APPROACH

The design was heavily influenced by principles set out by the FRP and dataflow models.

3.1 Reactive

The system is *reactive*, as close as possible to the definition of the Reactive Manifesto [12].

The system is *responsive*, meaning it should be able to handle time-sensitive scenarios if at all possible. This is the cornerstone of usability and utility, but more than that, it enables quick error-detection and error-handling.

The system is *resilient*, meaning it is able to recover robustly and gracefully after a failure, due to the fact that nodes in the dataflow graph are completely independent and recovery of each one can be done in isolation. Another thing to note here is that special error messages are built-in and make it very easy to propagate errors between *components*, in case the error-handling part of a component is decoupled from the computational logic. This leads to much more robust architectures for large-scale systems, where fault-tolerance is mission-critical.

The system is *elastic*, meaning it will adjust itself depending on the available resources and demanded workload. For instance, the granularity of the graph (i.e. number of nodes) is adjusted so as to match a heuristic-based value (e.g. total number of threads).

The system is *message-driven*, meaning it relies solely on asynchronous message-passing for inter-component communication leading to loose coupling, isolation, location transparency and the error propagation mentioned above. Location transparency is critical to preserve the semantics whether on a single host or a machine cluster.




Figure 4: Reactive properties

3.2 Heterogeneous

One of the major concerns while designing the framework was the ability to deploy it anywhere, from low-cost robots to mainframes. Obviously, such attribute would require a very flexible runtime environment. To satisfy this requirement, the strategy design pattern was used for evaluation, meaning that the core system only builds the internal representation of the dataflow graph and partitions it across the available computational resources. From there onwards, each partial graph can be evaluated by a different *EvaluationStrategy* (see Implementation chapter), which could interpret it using the Java 8 Streams library or even compile into CUDA code for execution on a GPU.

Figure 5 illustrates a simple example of a robot application pipeline, where input to the dataflow graph is what the robot's camera senses, and after some image processing and some computation-heavy decision making, a command to an actuator of the robot is executed. Orange nodes are deployed on the robot's on-board computer, the green node is deployed on an off-board GPU and the red node is deployed on the lab's main cluster.




Figure 5: Heterogeneity pipeline

3.3 Extensible

While doing my internship at NCSR "Demokritos", I realized that the goal I was pursuing was impossible to reach closure in a 6-month period by a humble undergraduate student. Therefore, I decided that everything should be implemented in a way that will allow future contribution by me or other researchers/developers.

With that concept in mind, I started generalizing and abstracting away everything I had done so far with the hope that the framework will raise attention later on. I can now say I am satisfied with the level of abstraction the core system has reached and I hope the stressful refactoring that the framework went through will blossom in the form of future contributions.

3.4 Abstract

The framework is *abstract* in terms of implementation details, as it is completely agnostic of any machine-specific or runtime-specific requirements. It is designed as a unifying conceptual base for further refined extensions and careful consideration was taken not to restrict in any aspect, architectural or not. The above was achieved by making many parts of the core system pluggable, which allows for easy refactoring on most of its internal functionality. Moreover, the internal graph representation does not include information on

how a node is executed, but only what are its semantics.

4. IMPLEMENTATION

This section examines the major characteristics of the framework's implementation.

4.1 Notifications

Every value passed through the framework's *streams* is wrapped inside a *Notification* object, which discriminated stream values into three categories: *onNext* (when the stream provides a regular value), *onError* (when an error occurs) and *onComplete* (when the stream completes its output).

4.2 The Reactive Streams Standard

In order to make the framework easy to integrate with other stream and/or dataflow technologies, I decided that every input/output node (i.e. publisher/subscriber in the PubSub terminology or source/sink in the Dataflow terminology) should implement the interfaces that the Reactive Streams Standard (RSS) define. RSS is an initiative to provide a standard for asynchronous stream processing with non-blocking back pressure. This encompasses efforts aimed at runtime environments (JVM and JavaScript) as well as network protocols [13].

A sink node (output) should implement the Subscriber interface ¹, which essentially defines three methods corresponding to reactions to a *Notification*, one for each of the categories mentioned above.

A source node (input) should implement the Publisher interface ², which defines a single method *subscribe(Subscriber s)*, where a Subscriber requests the Publisher to start emitting values.

Many existing technologies provide these interfaces, or at least adapters from their internal representations, and therefore they are very easy to be integrated to the framework.

¹<http://www.reactive-streams.org/reactive-streams-1.0.0-javadoc/org/reactivestreams/Subscriber.html>

²<http://www.reactive-streams.org/reactive-streams-1.0.0-javadoc/org/reactivestreams/Publisher.html>

4.3 Software structure

The core system ³ is organized in the following top-level packages:

<i>org.rhea_core.internal</i>	all internal functionalities such as graphs, expressions and notifications
<i>org.rhea_core.evaluation</i>	everything associated with the evaluation of the constructed dataflow graph
<i>org.rhea_core.distribution</i>	everything associated with distributing the evaluation across the available computational resources
<i>org.rhea_core.optimization</i>	includes some built-in optimizers for adjusting the granularity of the graph
<i>org.rhea_core.io</i>	defines the interfaces that sources/sinks should implement
<i>org.rhea_core.util</i>	helpful utilities needed throughout the project

4.4 Internal representation

For representing the internal structure of the dataflow graph, the *JGrapht* open-source Java library was used, which provides many graph data structures and common graph-theory algorithms ⁴. The main class representing the internal dataflow graph is *FlowGraph*, which is located in the *org.rhea_core.internal* package.

4.5 Stream variables


The construction of the aforementioned internal graph is always implicit, through a rich set of stream operators. The only data type handled by the programmer is *org.rhea_core.Stream*, which contains the parametric type of its single output. Each *Stream* object contains internally a *FlowGraph*, which is only to be accessed and manipulated by the internal module, evaluation strategies and optimizers.

Figure 6 illustrates the dataflow graph(left) produced by the framework code(right) using the *Stream* data type.

These variables can be used, preferably together with their parametric type, and reused in different parts of the graph. This is necessary, for instance, when you wish to split a node's output to different inputs. Figure 7 shows such an example.


³<https://github.com/rhea-flow/rhea-core>

⁴<http://jgrapht.org/>



```
Stream.zip(
    Stream.range(1, 10),
    Stream.range(1, 10),
    (x, y) -> x + y);
```

Figure 6: Simple example



```
Stream<Int> st = Stream.range(1, 10);
st.map(f).print();
st.filter(g).print();
```

Figure 7: Split example

4.6 Stream operators

4.6.1 Creation

4.6.2 Combining

4.6.3 Filtering

4.6.4 Conditional

4.6.5 Transformational

4.6.6 Feedback

4.6.7 Error-handling

4.6.8 Backpressure

4.6.9 Utility

4.7 Evaluation

4.8 Distribution

4.8.1 Hazelcast

Due to the RSS being in its infant stage, no working implementation exists for superimposing a network protocol onto it (e.g. RSS over TCP).

4.8.2 Machine configuration

4.9 Optimization

4.9.1 Heuristic-based graph transformations

4.9.2 Network-aware node placement

4.10 Serialization

4.10.1 The Serializer Interface

4.10.2 JsonIO

5. APPLICATIONS

5.1 Hamming numbers

5.2 Camera surveillance

5.3 Robot control panel

5.4 Robot hospital guide

6. RELATED WORK

6.1 Big Data

6.1.1 GoogleDataflow

6.1.2 TensorFlow

6.1.3 Akka

6.1.4 dispel4py

6.1.5 Spark

6.1.6 Naiad

6.2 Robotics

6.2.1 Flowstone

6.2.2 Yampa

6.2.3 roshask

6.3 Internet of Things

6.3.1 NodeRed

6.3.2 NoFlo

7. FUTURE WORK

- 7.1 More evaluation strategies**
- 7.2 Dynamic reconfiguration**
- 7.3 Advanced network profiling**
- 7.4 DSLs for framework extension**
- 7.5 Advanced fault-tolerance**
- 7.6 Integration with other technologies**
- 7.7 Stream reasoning**

8. CONCLUSIONS

ABBREVIATIONS, INITIALS AND ACRONYMS

A table of all abbreviations used throughout the thesis follows.

FRP	Functional Reactive Programming
JVM	Java Virtual Machine
NCSR	National Centre for Scientific Research
ROS	Robot Operating System
IoT	Internet of Things
CPU	Central Processing Unit
TCP	Transmission Control Protocol
PubSub	Publish/Subscribe
OOP	Object-oriented Programming
UML	Unified Modelling Language
GPU	Graphics Processing Unit
DSL	Domain-specific Language
RSS	Reactive Streams Standard

REFERENCES

- [1] C. Elliott and P. Hudak, “Functional reactive animation,” in *ACM SIGPLAN Notices*, vol. 32, pp. 263–273, ACM, 1997.
- [2] E. Amsden, “A survey of functional reactive programming,” *Unpublished*, 2011.
- [3] Z. Wan, W. Taha, and P. Hudak, “Real-time frp,” in *ACM SIGPLAN Notices*, vol. 36, pp. 146–156, ACM, 2001.
- [4] C. M. Elliott, “Push-pull functional reactive programming,” in *Proceedings of the 2nd ACM SIGPLAN symposium on Haskell*, pp. 25–36, ACM, 2009.
- [5] P. Hudak, A. Courtney, H. Nilsson, and J. Peterson, “Arrows, robots, and functional reactive programming,” in *Advanced Functional Programming*, pp. 159–187, Springer, 2003.
- [6] J. Peterson, P. Hudak, and C. Elliott, “Lambda in motion: Controlling robots with haskell,” in *Practical Aspects of Declarative Languages*, pp. 91–105, Springer, 1999.
- [7] Z. Wan, W. Taha, and P. Hudak, “Event-driven frp,” in *Practical Aspects of Declarative Languages*, pp. 155–172, Springer, 2002.
- [8] J. Hughes, “Generalising monads to arrows,” *Science of Computer Programming*, vol. 37, no. 1–3, pp. 67 – 111, 2000.
- [9] R. Paterson, “A new notation for arrows,” *SIGPLAN Not.*, vol. 36, pp. 229–240, Oct. 2001.
- [10] M. Quigley, K. Conley, B. Gerkey, J. Faust, T. Foote, J. Leibs, R. Wheeler, and A. Y. Ng, “Ros: an open-source robot operating system,” in *ICRA workshop on open source software*, vol. 3, p. 5, 2009.
- [11] E. Gamma, *Design patterns: elements of reusable object-oriented software*. Pearson Education India, 1995.
- [12] B. Jonas, D. Farley, R. Kuhn, and M. Thompson, “Reactive manifesto.” <http://www.reactivemanifesto.org/>, 2014.
- [13] “Reactive streams standard.” <http://www.reactive-streams.org/>, 2015.